



KUBERNATIC

Cloud Native Rejekts [EU'22]

**Using defaults for Deployments?  
Is it safe and sound?**





## Koray Oksay

Site Reliability Engineer

✉ koray@kubermatic.com

🐦 @korayoksay

👤 kubermatic/kubermatic

## What Else?

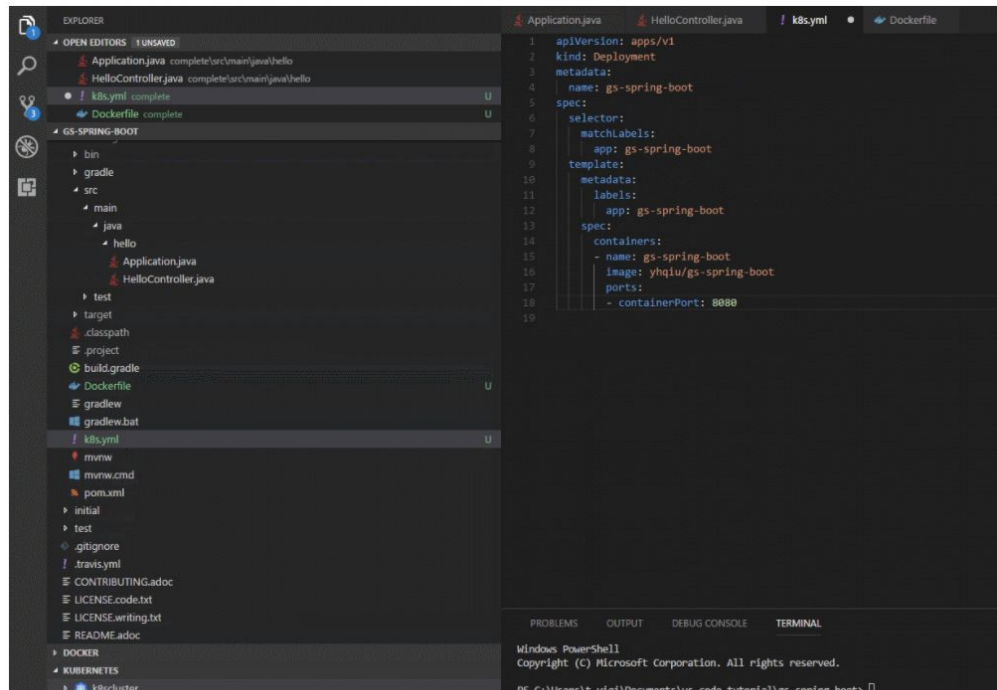
- Consultant / Site Reliability Engineer
- Working remotely from Istanbul
- Interested in Linux, Kubernetes, and cloud technologies
- CKA | CKAD | CKS | RHCE

# Kubernetes Deployment

- [kubernetes.io/docs](https://kubernetes.io/docs)
- VS Code Kubernetes plugin

```
controllers/nginx-deployment.yaml

apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```



## Quality of Service Classes

- **Guaranteed:**
  - Top priority
  - Pods will not be killed until they exceed their limits
- **Burstable:**
  - Minimal resource guarantee
  - Pods will not be killed until they exceed their limits and not any best-effort pods exist.
- **Best Effort:**
  - **Lowest priority**
  - **First to kill when the system is out of memory**

# Add Resource Requests and Limits

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.14.2
        ports:
        - containerPort: 8080
        resources:
          requests:
            memory: "64Mi"
            cpu: "250m"
          limits:
            memory: "128Mi"
            cpu: "500m"
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.14.2
        ports:
        - containerPort: 8080
        resources:
          requests:
            memory: "128Mi"
            cpu: "500m"
          limits:
            memory: "128Mi"
            cpu: "500m"
```

# Define Probes

- Readiness:
  - Checks if the container is ready to accept traffic
- Liveness:
  - Checks if the application is running and makes progress

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 8080
            readinessProbe:
              exec:
                command:
                  - cat
                  - /tmp/healthy
            livenessProbe:
              httpGet:
                path: /liveness
                port: 8080
```

# Add Security Context

- Make sure you run your container as non-root user (Container escape?).
- Have read-only root filesystem (do not change filesystem of your container).
- Process should not gain more privileges than its parent.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 8080
          securityContext:
            runAsUser: 1000
            runAsNonRoot: true
            allowPrivilegeEscalation: false
            readOnlyRootFilesystem: true
```

# Terminate Your Pods Gracefully

The steps to terminate a pod:

- Pod state is changed to **Terminating** and **preStop** hook is executed (if defined)
- Send **SIGTERM** to PID 1 of all containers (This might be tricky!)
- Wait for 30 seconds (Default value for **TerminationGracePeriodSeconds**)
- Send **SIGKILL**

Tricky part: if your pods are always wait for 30 seconds (or other defined value)

- Check your Dockerfile for the CMD instruction:
  - **CMD ./app.py** → This starts a shell and runs the app (/bin/sh -c ./app.py)
  - **CMD [ "./app.py" ]** → This runs the app without a shell, so it gets **PID 1**



# Distribute Your Pods

- By default, kube-scheduler runs the pods on any available node
- This could cause all your 3 replicas running on the same node!
- Ask the scheduler to run your pods on different nodes - a.k.a. **AntiAffinity**

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 8080
          terminationGracePeriodSeconds: 60
      affinity:
        podAntiAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            - labelSelector:
                matchExpressions:
                  - key: app
                    operator: In
                    values:
                      - nginx
              topologyKey: kubernetes.io/hostname
```

# All Combined

```
...
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 8080
    resources:
      requests:
        memory: "64Mi"
        cpu: "250m"
      limits:
        memory: "128Mi"
        cpu: "500m"
    readinessProbe:
      exec:
        command:
        - cat
        - /tmp/healthy
    livenessProbe:
      httpGet:
        path: /liveness
        port: 8080
    securityContext:
      runAsUser: 1000
      runAsNonRoot: true
      allowPrivilegeEscalation: false
      readOnlyRootFilesystem: true
    terminationGracePeriodSeconds: 60
  affinity:
    podAntiAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
      - labelSelector:
          matchExpressions:
          - key: app
            operator: In
            values:
            - nginx
        topologyKey: kubernetes.io/hostname
```

# Thank you!

✉ [koray@kubermatic.com](mailto:koray@kubermatic.com)

🐦 [@kubermatic](https://twitter.com/kubermatic)

🐙 [kubermatic/kubermatic](https://github.com/kubermatic/kubermatic)